

DSC Analysis

Technical Guide

Distance Sample Count — Wildlife Survey Analysis Pipeline

Generated August 14, 2026

Workflow id: **dsc_analysis**

1. Overview

The **dsc_analysis** workflow ingests Distance Sample Count (DSC) wildlife survey data from EarthRanger and produces structured, analysis-ready datasets for use in population density modelling. The workflow supports **multiple surveys** in a single run, each defined by its own EarthRanger connection, patrol type, survey time window, and transect spatial group. Each survey is further split automatically into its **detected activity periods**, so a site visited in e.g. both February and June produces two independent, period-tagged output sets rather than one combined one.

For each survey period the workflow delivers:

- A **metadata CSV** — survey-level observations including transect IDs, observer counts, team members, and event types
- A **field effort CSV** — per-team-member daily distance travelled, duration, and man-hours, derived from EarthRanger subject GPS tracks
- An **analysis data CSV** — wildlife observation events enriched with off-transect distances, orthogonal distances, estimated animal positions, and satellite-derived environmental covariates (NDVI, terrain slope)
- An **events GeoPackage** — spatial point layer of wildlife observations with key distance sampling geometry fields
- A **transect areas GeoPackage** — visited transect corridors (buffered) labelled with mean NDVI, slope, and PALSAR woody cover values from Google Earth Engine
- A **transect lines GeoPackage** — visited transect centrelines (unbuffered), reprojected to EPSG:4326

Output summary

Output type	Format	Description
{survey}_{period}_analysis_metadata	CSV	Survey event metadata: transect IDs, team members, observer counts, event types
{survey}_{period}_field_effort	CSV	Per-team-member daily distance, duration, and man-hours from GPS tracks
{survey}_{period}_analysis_data	CSV	Wildlife observations with distance fields, animal position estimates, NDVI, and slope covariates
{survey}_{period}_events	GeoPackage	Spatial point layer of wildlife observation events
{survey}_{period}_transect_areas	GeoPackage	Visited transect corridors (buffered) with NDVI, slope, and woody cover labels (EPSG:4326)
{survey}_{period}_transect_lines	GeoPackage	Visited transect centrelines (unbuffered) in EPSG:4326

Note: {survey} is the survey name as defined in the connection configuration and {period} is the detected activity period label (e.g. 2026_02). All six output file sets are produced independently for each survey period.

2. Dependencies

2.1 Python packages

The workflow declares six versioned packages from the Ecoscope prefix.dev channels:

Package	Version	Channel
ecoscope-platform	>=2.15.0, <2.16.0	ecoscope-workflows
ecoscope-workflows-ext-custom	0.1.0rc14.*	ecoscope-workflows-custom
ecoscope-workflows-ext-ste	0.0.0rc1.*	ecoscope-workflows-custom
ecoscope-workflows-ext-distance-sample-counts	1.0.11.*	ecoscope-workflows-custom
pydeck	0.9.2	conda-forge
opentelemetry-sdk	>=1.20.0, <2.0.0	conda-forge

Note: Earlier revisions of this workflow additionally pinned separate ecoscope-workflows-core / ecoscope-workflows-ext-ecoscope packages plus site-specific extensions (ecoscope-workflows-ext-big-life, -ext-mnc, -ext-ate). These have been consolidated into the single ecoscope-platform package and removed where no longer required.

2.2 Google Earth Engine connection

A Google Earth Engine (GEE) service account connection is required (**set_gee_connection**). The GEE client is used to build satellite imagery composites and label transect lines with environmental covariates. The same GEE client is shared across all surveys in a single run.

2.3 EarthRanger connections

One EarthRanger connection is required *per survey*. Each entry in **connection_config** pairs an EarthRanger server and patrol type ID with one or more survey definitions (survey name, time window, and transect spatial group ID). Connections are split and distributed to parallel per-survey pipelines via **split_connection_configs**.

2.4 EarthRanger data requirements

The workflow expects the following event types to be present in EarthRanger for each patrol:

Event type	Role
distancecountpatrol_rep	Survey metadata event — carries transect ID, team members, and observer count
distance_count_patrol_metadata	Alternative metadata event type (also retained during metadata filtering)
distancecountwildlife_rep	Wildlife observation event — carries species, total count, distance to centre, radial angle, and juvenile count

Note: Field effort computation additionally requires that team members recorded on patrol metadata events (Team Members / reported_by) match named subjects/trackers in EarthRanger, so their GPS observations can be retrieved and converted into daily distance and duration statistics.

3. Input Configuration

3.1 Workflow-level parameters

Parameter	Description
workflow_details	Human-readable name and optional description for this workflow run — used for dashboard registration
time_range	Required on all Ecoscope workflows. Used for timestamp display and UTC conversion only — does not filter which patrol events are fetched. Set broadly to cover all surveys in the run. Each survey's own time window controls data fetching.
gee_client	Google Earth Engine service account connection
connection_config	Array of EarthRanger connection entries, one per survey (see Section 3.2)

3.2 Connection config entry (per survey)

Each entry in **connection_config** defines one EarthRanger site and one or more surveys to run against it:

Field	Type	Description
earthranger.server	Connection	EarthRanger data source connection
earthranger.patrol_type_id	String	Numeric or UUID identifier for the patrol type containing DSC surveys
surveys[].surveyName	String	Unique name for this survey (used as output filename prefix)
surveys[].time_range	Object	Survey-specific data fetch window. Patrol events within this since / until range are retrieved from EarthRanger for this survey. This is the field that controls which data is pulled — not the top-level time_range.
surveys[].erSpatialTransectsGroupId	String	EarthRanger spatial group ID that contains the transect line features

4. Data Ingestion

4.1 Connection splitting

The combined **connection_config** is split into individual per-survey connection objects via **split_connection_configs**. All subsequent steps that require per-survey data use **mapvalues** to fan out over this list in parallel.

4.2 Patrol events

For each survey, **fetch_patrol_events** retrieves all patrols matching the configured patrol type ID within the survey time window. The resulting patrol DataFrame is indexed on the **id** column via **set_dataframe_index** to enable subsequent join operations.

4.3 Survey period splitting

Before any further processing, each survey's patrol events are passed through **split_survey_by_period**, which detects the distinct activity periods within the fetched patrol events (e.g. a site visited in February and again in June) and splits them into independent, period-keyed branches. **unpack_period_connection_survey** and **unpack_period_patrol_events** then extract the connection/survey config and patrol events for each period, and **get_survey_period_label** derives the period label (e.g. **2026_02**) used in output filenames. From this point on, every downstream step fans out over survey *periods* rather than surveys.

4.4 Patrol transects

Transect lines are fetched per survey period via **fetch_transects**, using the spatial group ID specified in the survey's connection config. Because a survey's transect geometry does not change between periods, the same transects are fetched once per period (duplicated across periods of the same survey) so that each period's output is self-contained. The transects are GeoDataFrames in EPSG:4326 as returned by EarthRanger.

4.5 Survey observation events

Individual wildlife observation events are fetched from the patrol event IDs produced by period splitting using **fetch_events** (**chunk_size**: 75). This two-step approach — fetch patrols first, then fetch the events within them — is required because the EarthRanger API does not support direct patrol-type filtering on the events endpoint.

4.6 EarthRanger server name

The EarthRanger server name (site identifier) is retrieved per survey period via **get_server_name** and zipped with the survey event DataFrame. It is passed to **process_events_details** as the **client** argument so that EarthRanger field IDs can be resolved to human-readable display names. The domain name is retrieved separately via a second **get_server_name** call for use in filename/traceability construction.

5. Event Processing Pipeline

5.1 Metadata event branch

Survey metadata events (**distancecountpatrol_rep** and **distance_count_patrol_metadata**) are processed through a dedicated branch to produce the *analysis metadata* output:

Step	Task	Purpose
1	filter_row_values	Retain only metadata event types from the raw event fetch
2	parse_df_point	Extract latitude and longitude columns from the GeoDataFrame geometry
3	normalize_json_column (reported_by)	Flatten the reported_by JSON field into flat columns
4	join_list_column (reported_by__name)	Collapse list values in reported_by__name to a comma-separated string
5	process_events_details	Resolve EarthRanger field IDs to display names (map_to_titles: true, ordered: true)
6	normalize_json_column (event_details)	Flatten the event_details JSON into flat columns
7	map_columns	Retain and rename key columns: time, event_type, serial_number, latitude, longitude, reported_by, Transect ID, Team Members, Number of Observers
8	parse_list_column / join_list_column (Team Members)	Parse and flatten Team Members list into a comma-separated string
9	drop_null_columns	Drop any columns that are entirely null after normalisation

The resulting DataFrame is persisted as **{survey_name}_{period}_analysis_metadata.csv**.

5.2 Wildlife observation event branch

Wildlife observation events (**distancecountwildlife_rep**) pass through a separate normalisation branch before being merged with patrol-level metadata:

Step	Task	Purpose
1	select_columns (event_details)	Extract only the event_details column from the raw event fetch for merging with the patrol DataFrame
2	merge_two_dataframes (left join on index)	Join the patrol DataFrame with the event_details column using the patrol id index
3	convert_column_timezone	Convert the time column to UTC using the timezone from the global time_range (used for display/conversion only, not for data filtering)
4	normalize_json_column (event_details)	Flatten the merged event_details JSON into flat columns

Step	Task	Purpose
5	map_columns (patrol rename)	Rename flattened distance sampling fields to standardised names (see table below)
6	format_text_column (transect_id, lower)	Lowercase transect_id so it matches the lowercased transect names used downstream for name-based transect matching
7	bfill_within_patrols / ffill_within_patrols	Backward- then forward-fill transect_id and num_observers within each patrol (group_col: patrol_serial_number) to propagate metadata event values to wildlife observation rows
8	filter_row_values (wildlife sightings + patrol "conducted" markers)	Retain distancecountwildlife_rep, distancecountpatrol_rep, and distance_count_patrol_metadata after fill propagation. The two patrol marker types are kept alongside sightings so that a transect with a patrol but zero sightings still has a row to test against the transect corridor downstream (Section 6.3) — without this, that transect would never be flagged as visited
9	add_constant_column (survey_id)	Stamp each row with the survey+period base name (e.g. olaremotorogi_2026_02) as a survey_id identifier

Note: A later step, filter_row_values (distancecountwildlife_rep only) — Section 6.3 — narrows back down to sightings-only immediately before the merge that feeds analysis_data.csv / events.gpkg. The patrol marker rows are only ever needed to establish which transects were visited; they do not appear in either of those two exports.

5.3 Field name mapping

The **map_columns** step in the wildlife observation branch renames the flattened EarthRanger field codes to analysis-ready column names:

Source field (EarthRanger)	Target column
event_details__distancecountpatrol_numberofobservers	num_observers
event_details__distancecountpatrol_teammembers	Team Members
event_details__distancecountpatrol_transectid	transect_id
event_details__distancecountwildlife_distancetocentre	dist_to_centre
event_details__distancecountwildlife_numberofjuveniles	num_juveniles
event_details__distancecountwildlife_radialangle	radialangle
event_details__distancecountwildlife_species	species
event_details__distancecountwildlife_totalcount	totalcount
event_details__Transect_ID (alt. form, underscore)	transect_id_v2
event_details__Transect ID (alt. form, space)	transect_id_v2
event_details__Team_members (alt. form)	Team Members

Source field (EarthRanger)	Target column
event_details__Number_of_observers (alt. form)	num_observers

Note: The mapping includes both the snake_case EarthRanger internal names and alternative title-case forms to handle variation across EarthRanger server configurations — including servers that render the field label with a literal space ("Transect ID") rather than an underscore. `raise_if_not_found` is set to false so that missing columns are silently ignored.

Note: Both alt-form transect ID fields land in `transect_id_v2`, a separate column from `transect_id` (populated by the `distancecountpatrol_transectid` mapping above). A following `map_columns` step ("Normalize transect_id column name across schemas") renames `transect_id_v2` → `transect_id` so both schemas converge on one column before the `parse/lowercase/fill` steps below run.

5.4 Field effort computation

In parallel with the metadata and wildlife observation branches, **compute_field_effort** derives a daily field-effort summary per team member from the metadata event DataFrame (post `drop_null_columns`) and the survey's EarthRanger connection:

Step	Description
1. Roster reconstruction	Team Members (or <code>reported_by</code> as a fallback) is parsed per event and grouped by recorder and <code>survey_date</code> to build a daily team roster. If no member names are recorded, a placeholder roster of size 3 (<code>DEFAULT_TEAM_SIZE</code>) is generated from the recorder so that a team size estimate is still produced.
2. Subject resolution	Roster member names are matched against EarthRanger subjects (<code>get_subjects</code>) to resolve subject IDs. If no match is found for individual members, the recorder name is matched against subjects as a fallback.
3. Observation window	Each <code>survey_date</code> is assigned a fixed daily observation window (04:00–13:00 UTC) used to bound the GPS observation fetch for that date.
4. Track retrieval	Subject GPS observations are fetched per <code>survey_date</code> via <code>get_subject_observations</code> and converted to relocations, then to trajectories (<code>process_relocations / relocations_to_trajectory</code>), filtered to segments between 0.1 m and 100 km, 1 second and 6 hours, and 1–100 km/h.
5. Distance/duration summary	Trajectory segments are summarised per subject/day (<code>summarize_df</code>) into total distance (km) and total duration (h), then merged back onto the roster.
6. Man-hours	<code>man_hours</code> is computed as <code>team_size × duration</code> for each subject/day row.

The resulting DataFrame — columns `survey_date`, `id`, `name`, `team_size`, `distance`, `duration`, `man_hours` — is persisted as **{survey_name}_{period}_field_effort.csv**.

Note: If no GPS observations are returned for a survey period, `compute_field_effort` returns an empty DataFrame with the expected columns rather than raising an error, so downstream persistence is skipped gracefully for that period.

6. Transect Processing

6.1 CRS conversion to UTM

Distance sampling calculations require a metric coordinate reference system. The UTM zone is estimated automatically from the transect geometry via **estimate_utm_crs**. Both the transect GeoDataFrame and the wildlife observation GeoDataFrame are then reprojected to this UTM CRS via **reproject_gdf** before any distance calculations are performed.

6.2 Transect simplification and buffering

Before spatial intersection tests, the transect lines are prepared as follows:

Step	Task	Parameters	Purpose
1	merge_transect_lines	—	Merge all transect line segments for the survey into a unified GeoDataFrame keyed by transect name
2	simplify_transects	tolerance: 50 m	Reduce transect vertex count while preserving shape — improves intersection performance
3	buffer_transects	distance: 500 m, cap_style: round, single_sided: false, resolution: 5	Create a 500 m bilateral corridor around each transect for event intersection testing

Note: cap_style is round rather than flat. A flat cap terminates the corridor in a straight edge exactly at each transect's start/end point, which excluded events recorded just before reaching or just after finishing the transect even when they were only a few metres away — diagnostics showed the excluded events consistently projected to position 0 or line-length on the transect (i.e. clamped to an endpoint), not because they were actually far from it. The round cap extends the corridor in an arc past each endpoint instead, capturing those near-endpoint events.

6.3 Event–transect intersection filter

Wildlife observation events are tested against the buffered transect corridors via **flag_events_intersecting_transect**. This step adds an **intersects_transect** boolean column and records which transect each event falls within (transect_id_column: transect_id, transect_name_column: name, index_column: id).

Only transects visited by at least one intersecting event are retained for the final output via **filter_visited_transects**. This removes transects that were planned but not walked during the survey period. Because the upstream filter (Section 5.2, step 8) now keeps patrol "conducted" marker events alongside wildlife sightings, a transect that was walked but had zero sightings still has an intersecting event here and is retained as visited — it would previously have been dropped.

The patrol marker rows carried through **flag_events_intersecting_transect** exist only to establish visited status and do not belong in the wildlife observation exports. Downstream of transect labeling, a dedicated **filter_row_values** step ("Keep only wildlife sightings for analysis_data / events export", id: **filter_wild_analysis**) re-filters the same **flag_events_intersecting_transect** output back down to **distancecountwildlife_rep** only, before it is paired with the labelled transects and merged (Section 9). This keeps analysis_data.csv and events.gpkg wildlife-sightings-only while still letting patrol-only transects surface in transect_areas.gpkg / transect_lines.gpkg.

6.4 Re-projection to EPSG:4326 for export

After intersection filtering, visited transects are reprojected back to EPSG:4326 (**reproject_gdf**, `target_crs: "epsg:4326"`) before satellite imagery labeling and GeoPackage export.

Note: The visited, buffered corridor geometry from this step is exported as the `transect_areas` GeoPackage (Section 9). A second, independent branch filters the pre-buffer simplified centrelines (from Section 6.2, step 2) down to the same visited transect names (`filter_transect_lines_by_visited`) and reprojects them to EPSG:4326 for export as the `transect_lines` GeoPackage — the same visited subset, but with line rather than polygon geometry.

7. Geometric Calculations

7.1 Off-transect distance (observer position)

The perpendicular distance from the *observer's recorded GPS position* to the nearest transect centreline is calculated via **add_off_transect_distance** and stored in the column **off_transect_dist**. Events where the computed distance equals **-1** (indicating that no matching transect was found) are removed via **filter_rows** (op: "ne", value: -1).

7.2 Estimated animal position

The true animal position is estimated from the observer position, the recorded radial angle (**radialangle**), and the recorded distance from observer to detected animal (**dist_to_centre**) via **estimate_animal_positions**. The original observer geometry is preserved first by **add_orig_geometry** into the column **orig_geometry**; the main geometry column is then replaced with the estimated animal position.

7.3 Orthogonal distance (animal position)

After the animal position is estimated, the perpendicular distance from the *estimated animal position* to the transect centreline is computed via a second call to **add_off_transect_distance** and stored as **ortho_dist**. This is the key detection distance used in distance sampling density estimation models.

Events with **ortho_dist == -1** are again removed via **filter_rows** before the data proceeds to imagery labeling.

7.4 Distance calculation summary

Column	Task	Geometry used	Meaning
off_transect_dist	add_off_transect_distance (1st call)	Observer GPS position	Distance from observer to transect centreline — QA check
orig_geometry	add_orig_geometry	Observer GPS position	Preserved original observer geometry before position estimation
geometry	estimate_animal_positions	Computed radialangle + dist_to_centre	Estimated true animal location
ortho_dist	add_off_transect_distance (2nd call)	Estimated animal position	Perpendicular distance from animal to transect — used in DSC models

8. Satellite Imagery Labeling

Each transect line is labelled with two satellite-derived environmental covariates using Google Earth Engine. These covariates are commonly used as predictors in distance sampling detection function models.

8.1 HLS NDVI

A Harmonized Landsat Sentinel-2 (HLS) NDVI composite is built per survey via **build_hls_ndvi_image**:

Parameter	Value	Notes
ndvi_window_days	null	No fixed window — uses the full image archive available up to the survey start date
max_cloud_cover	30 %	Scenes with more than 30 % cloud cover are excluded from the composite
ndvi_band_name	NDVI_HSL	Output band name stored in the labelled transect column

The NDVI image is evaluated over each transect's geometry and the mean pixel value within each transect is computed via **label_features_with_image_stat** (scale: 30 m, reducer_key: mean, out_column: NDVI_HSL). The survey start date is added as a **img_date_hsl_ndvi** column for traceability.

8.2 Terrain slope

A slope layer is derived from the USGS SRTM 1-arc-second DEM (**USGS/SRTMGL1_003**) via **build_slope_image**. The mean slope value within each transect is then computed via a second call to **label_features_with_image_stat** (scale: 30 m, reducer_key: mean, out_column: slope).

8.3 PALSAR woody cover

A woody cover layer is built from PALSAR imagery via **build_palsar_woody_cover_image**, using a dry-season window of **July–September** (dry_season_start_month: 7, dry_season_end_month: 9) and output band name **WoodyCover**. Like the HLS NDVI composite, this image is built per survey period, anchored to that period's minimum patrol event date (zip_aoi_min_date).

The mean woody cover value within each transect is computed via a third call to **label_features_with_image_stat** (scale: 50 m, reducer_key: mean, out_column: WoodyCover), chained after the slope-labelled transects via **zip_woody_cover_image**.

8.4 Imagery labeling workflow

The transects are converted to a Google Earth Engine FeatureCollection via **to_ee_feature_collection** before labeling. The period's minimum patrol event date (**get_survey_min_date**) is both added as a column to the transect DataFrame and passed to the NDVI image builder as the **since** parameter to anchor the composite date.

8.5 Final transect column selection

After labeling, transects are trimmed to the columns required for merging and export:

Column	In <code>transect_areas.gpkg</code>	Included in merge	Description
<code>name</code>	Yes	Yes	Transect identifier (matches <code>transect_id</code> in patrol events)
<code>img_date_hsl_ndvi</code>	Yes	Yes	Period min. patrol date used as NDVI image anchor
<code>NDVI_HSL</code>	Yes	Yes	Mean HLS NDVI value along transect
<code>slope</code>	Yes	Yes	Mean terrain slope along transect
<code>WoodyCover</code>	Yes	Yes	Mean PALSAR woody cover value along transect
<code>geometry</code>	Yes	No	Transect corridor geometry (excluded from CSV merge)

Note: The separately exported `transect_lines.gpkg` (Section 6.4) carries only the transect name and unbuffered line geometry — it is not labelled with `NDVI_HSL`, `slope`, `WoodyCover`, or `img_date_hsl_ndvi`.

Note: "Included in merge" means the column survives into `merge_filtered_patrols`' right-hand input (`exclude_geom`). `WoodyCover` does reach that merge step, but as of this revision it is then dropped by the final `analysis_data.csv` column selection (`select_patrol_event_cols`), which was not updated to include it alongside `NDVI_HSL` and `slope` — see Section 9 for the current `analysis_data.csv` column list.

9. Output Files

All outputs are written under **\$ECOSCOPE_WORKFLOWS_RESULTS**, into a dedicated subfolder per output type rather than the flat results root, so reviewers see files pre-sorted by type. Each subfolder is constructed once per run via **build_output_subfolder** (root_path: \$ECOSCOPE_WORKFLOWS_RESULTS, subfolder: <name>) and its return value is used as the root_path for the corresponding persist_df step. Six file sets are produced for each survey period. {survey} is the survey name defined in the connection config and {period} is the detected activity period label (e.g. 2026_02).

Subfolder	File	Format	Description
ER_Metadata/	{survey}_{period}_analysis_meta_data.csv	CSV	Survey metadata events: transect IDs, team members, observer counts, event types, lat/lon
ER_FieldEffort/	{survey}_{period}_field_effort.csv	CSV	Per-team-member daily field effort: survey_date, id, name, team_size, distance (km), duration (h), man_hours
ER_AnalysisData/	{survey}_{period}_analysis_data.csv	CSV	Wildlife observation events with all distance sampling fields: species, totalcount, num_juveniles, dist_to_centre, radialangle, off_transect_dist, ortho_dist, survey_id, orig_geometry (WKT), estimated geometry (WKT), NDVI_HSL, slope, img_date_hsl_ndvi, intersects_transect, transect_id, num_observers, time, serial_number, patrol_id, patrol_serial_number
ER_Events/	{survey}_{period}_events.gpkg	GeoPackage	Spatial point layer of wildlife observations. Columns: serial_number, transect_id, dist_to_centre, ortho_dist, intersects_transect, geometry
ER_SurveyAreas/	{survey}_{period}_transect_areas.gpkg	GeoPackage	Visited transect corridors — buffered polygons (EPSG:4326) with environmental covariates: name, img_date_hsl_ndvi, NDVI_HSL, slope, WoodyCover, geometry
ER_MasterTransects/	{survey}_{period}_transect_lines.gpkg	GeoPackage	Visited transect centrelines — unbuffered lines (EPSG:4326), the same visited subset as transect_areas but without environmental covariates

Note: The _events GeoPackage contains a subset of columns optimised for spatial QA and GIS workflows. The _analysis_data CSV contains the full column set including all covariates and is the primary input for distance sampling density estimation models. If a survey period has no visited transects or no field-effort GPS data, the corresponding file(s) for that period are simply not produced rather than raising an error.

Note: Visited transects now include those with a patrol "conducted" marker event and zero wildlife sightings (Section 6.3) — such transects appear in _transect_areas.gpkg / _transect_lines.gpkg but contribute no rows to _analysis_data.csv / _events.gpkg, which stay wildlife-sightings-only.

Note: WoodyCover is labelled onto transects and appears in `_transect_areas.gpkg` (Section 8.5), but as of this revision it is not included in the `_analysis_data.csv` column list above — only NDVI_HSL and slope are currently carried through to that file.

10. Workflow Execution Logic

10.1 Skip conditions

All tasks share a global default skip policy defined in **task-instance-defaults**:

- **any_is_empty_df** — skips the task if any upstream DataFrame dependency is empty
- **any_dependency_skipped** — skips the task if any upstream task was itself skipped

This means that if a survey returns no patrol events or no transects, all downstream tasks for that survey branch are skipped gracefully without raising an error. Many **groupbykey** steps additionally set **any_keyed_iterables_are_skips** (unpack_depth: 1) so that if either input iterable for a given key was itself a skip, that key's pairing is dropped instead of propagating a malformed pair downstream.

Note: The final publish-path aggregation step (Section 10.6 / dsc_publish handoff) deliberately omits **any_keyed_iterables_are_skips**: with many independent survey periods, it is normal for some individual period to have nothing persisted (e.g. no visited transects that period) while others succeed. **groupbykey**'s own **SkippedDependencyFallback** already drops just the skipped entries and keeps the rest, so using the stricter check there would discard every other period's real output.

10.2 Per-survey-period fan-out with mapvalues

Every data-bearing step is executed once per survey *period* via the **mapvalues** directive. The fan-out list originates from **split_connection_configs**, is then exploded per detected activity period by **split_survey_by_period**, and flows through all subsequent steps. Key fan-out points:

Step	mapvalues source
fetch_patrol_events	split_connection_config.return
period_split (activity period detection)	zip_conn_patrol_events.return
fetch_patrol_transects	period_connection_survey.return
fetch_events_from_ids	period_split.return
All event processing steps	Chained from fetch_events_from_ids
compute_field_effort	zip_connection_metadata.return
Transect CRS / distance steps	Chained from fetch_patrol_transects
GEE labeling steps	Chained from reproject_transects
All persist steps	Chained via zip_filename_* groupbykey

10.3 groupbykey — multi-input coordination

groupbykey (the successor to the older **zip_groupbykey** task) is used throughout the workflow to combine two or more per-period iterables into paired tuples before feeding them into a multi-argument **mapvalues** step. Key uses:

- **zip_conn_patrol_events** — pairs **split_connection_config** with **set_patrol_index** to feed **period_split**
- **zip_conn_survey_df** — pairs the event DataFrame with the EarthRanger server name to feed **process_events_details**

- **zip_connection_metadata** — pairs period_connection_survey with the metadata summary table to feed compute_field_effort
- **zip_patrol_transects** — pairs patrol events (UTM) with transects (UTM) to feed add_off_transect_distance
- **zip_aoi_min_date** — pairs the EE FeatureCollection with the period's minimum patrol date to feed build_hls_ndvi_image
- **zip_lines_visited** — pairs the pre-buffer simplified transect lines with the visited-transects list to feed filter_transect_lines_by_visited
- **zip_filename_*** — pairs dynamically constructed filenames with DataFrames to feed persist_df

10.4 Dynamic filename construction

Output filenames are constructed at runtime with **join_with_underscore**. The survey name and period label are first combined into a shared base name (**combine_survey_period_name**, e.g. olaremotorogi_2026_02) that is then joined with a fixed suffix per output:

Suffix	Output file purpose
analysis_metadata	Metadata events CSV
field_effort	Field effort summary CSV
analysis_data	Wildlife observations CSV
events	Events GeoPackage
transect_areas	Visited transect corridors (buffered) GeoPackage
transect_lines	Visited transect centrelines (unbuffered) GeoPackage

Note: The constructed filename is combined with the output's dedicated subfolder (Section 9) — built once per run via build_output_subfolder — at the persist_df step for that output, so each file lands in \$ECOSCOPE_WORKFLOWS_RESULTS/<subfolder>/<filename> rather than directly under \$ECOSCOPE_WORKFLOWS_RESULTS.

10.5 Fill propagation for metadata fields

Patrol events in EarthRanger are recorded sequentially: a metadata event (distancecountpatrol_rep) carrying the transect ID and observer count typically appears once per transect walk, while multiple wildlife observation events (distancecountwildlife_rep) follow it. Because both event types share the same patrol_serial_number, the workflow uses backward fill (**bfill_within_patrols**) followed by forward fill (**ffill_within_patrols**) to propagate the transect_id and num_observers values to every row within the patrol before filtering to wildlife observation events only. transect_id is lowercased immediately beforehand so it matches the lowercased transect names used for name-based transect matching.

10.6 Publish handoff (removed)

Earlier revisions of this workflow built a text widget (**create_text_widget_single_view**, title: "Files ready to publish") listing every persisted transects and patrol-events GeoPackage path across all surveys and periods, intended for a reviewer to copy into a companion **dsc_publish** workflow. That widget — along with its supporting **collect_patrol_events_paths** and **format_publish_paths** tasks — was never attached to the dashboard and has been removed as of this revision, superseded by the fully wired dashboard described in Section 11.

11. Interactive Dashboard

gather_dashboard assembles a dashboard of seven widgets. Every widget is built per survey period and then combined across periods by **merge_widget_views**, so the dashboard renders one set of tiles, map, and tables per period, switchable via the dashboard's grouper controls.

11.1 Groupers

set_groupers defines two facet keys, both consumed by **gather_dashboard**:

Index name	Facets by
group_id	The originating survey (one per connection_config entry)
period	The detected activity period within that survey (Section 4.3)

Note: Earlier revisions left groupers empty and shipped a dashboard with no widgets attached (Section 10.6). `group_id` and `period` mirror the same `{survey}_{period}` labelling used for output filenames (Section 9), so the dashboard's facet selector lines up one-to-one with the file naming a reviewer already sees in the exports.

11.2 Widgets

Widget	Type	Source
Total Distance Surveyed (km)	Stat tile	<code>dataframe_column_sum (distance, field_effort.return)</code> → <code>total_distance_grouped_widget</code>
Total Effort (hours)	Stat tile	<code>dataframe_column_sum (duration, field_effort.return)</code> → <code>total_effort_grouped_widget</code>
Number of Transects Surveyed	Stat tile	<code>dataframe_column_nunique (name, filter_areas_by_metadata.return)</code> → <code>total_transects_grouped_widget</code>
Survey Overview Map	Map	<code>draw_map (survey_overview_map)</code> → <code>survey_overview_map_widget</code>
Field Effort	Table	<code>draw_table (field_effort.return)</code> → <code>field_effort_table_widget</code>
Survey Metadata	Table	<code>draw_table (drop_null_cols.return)</code> → <code>metadata_grouped_table_widget</code>
Analysis Data	Table	<code>draw_table (select_analysis_table_cols.return)</code> → <code>analysis_data_table_widget</code>

Note: Total Transects counts distinct transect names in `filter_areas_by_metadata` — transects present in this period's metadata after the buffer/metadata-presence filter (Section 6.2 onward), which runs earlier in the pipeline than the final visited-transect filter (Section 6.3). The count on the dashboard tile can therefore differ slightly from the number of transects in that period's `transect_areas.gpkg` / `transect_lines.gpkg`.

11.3 Survey Overview Map

The map overlays three layers, each built with **ecoscope_workflows_ext_custom.tasks.results.create_geojson_layer** or **create_scatterplot_layer** and combined via **groupbykey (combined_survey_map_layers)**:

Layer	Source	Style
Transect Areas	format_transect_names.return (buffered, visited transects)	Green fill/outline, 55% opacity
Transect Lines	reproject_transect_lines.return (unbuffered centrelines)	Black outline, 55% opacity
Events	reproject_events_map.return (wildlife observations, reprojected to EPSG:4326)	Orange points, 75% opacity

Each layer's **skipif** includes **all_geometry_are_none** alongside the standard empty-dependency conditions, so a period with no visited transects, for example, simply omits that layer rather than failing the map. The view is auto-framed to the transect extent: **envelope_gdf** (expansion_factor: 1.35) computes a bounding envelope around the visited transects, and **compute_view_state_from_gdf** (pitch: 0, bearing: 0, max_zoom: 15) derives the pydeck view state from it. The combined layers and view state are zipped together (**zip_layers_with_viewstate**) and rendered by **draw_map** (max_zoom: 10, legend placement: bottom-right).

11.4 Tables

All three table widgets use **draw_table** with sorting, filtering, and download enabled (**table_config**: enable_sorting, enable_filtering, enable_download all true), rendered to HTML and persisted via **persist_text** before being wrapped as single-view table widgets:

Table	Source data	Notes
Field Effort	field_effort.return	Same DataFrame persisted as field_effort.csv (Section 5.4)
Survey Metadata	drop_null_cols.return	Same DataFrame persisted as analysis_metadata.csv (Section 5.1)
Analysis Data	select_analysis_table_cols.s.return	A dedicated column subset of select_patrol_event_cols.return (the analysis_data.csv source) with geometry columns dropped for in-dashboard display

Note: select_analysis_table_cols uses raise_on_missing: false, so a period missing one of the listed columns (e.g. other_species) still renders a table with the remaining columns rather than failing.

12. Software Versions

Package	Version pinned
ecoscope-platform	>=2.15.0, <2.16.0
ecoscope-workflows-ext-custom	0.1.0rc14.*
ecoscope-workflows-ext-ste	0.0.0rc1.*
ecoscope-workflows-ext-distance-sample-counts	1.0.11.*
pydeck	0.9.2
opentelemetry-sdk	>=1.20.0, <2.0.0

Note: All packages are resolved from the prefix.dev Ecoscope conda channels. The wildcard patch-version pin (.*) allows bug-fix releases to be picked up automatically while keeping minor and major versions locked.